

Where the agent wins: 1-node bench → 7-node bench → 7-node training (per-step contribution)

